

基於注意力機制多模態融合之
穩健駕駛狀態監控系統
A Robust Driver State Monitoring System Utilizing
Attention-Based Multi-Modal Fusion



研究生：張善為 (Shang-Wei Chang)
指導教授：許超雲 (Prof. Chau-Yun Hsu)
卓立 (Prof. Li Cho)

大同大學
電機工程研究所
碩士論文

Thesis for Master of Science
Department of Electrical Engineering
Tatung University

中華民國 一一五年八月
August 2026

基於注意力機制多模態融合之
穩健駕駛狀態監控系統
A Robust Driver State Monitoring System Utilizing
Attention-Based Multi-Modal Fusion



研究生：張善為 (Shang-Wei Chang)
指導教授：許超雲 (Prof. Chau-Yun Hsu)
卓立 (Prof. Li Cho)

大同大學
電機工程研究所
碩士論文

Thesis for Master of Science
Department of Electrical Engineering
Tatung University

中華民國 一一五年八月
August 2026

大 同 大 學
電 機 工 程 研 究 所
碩 士 學 位 論 文

中文標題

基於注意力機制多模態融合之穩健駕駛
狀態監控系統

ENGLISH TITLE

A Robust Driver State Monitoring System Utilizing
Attention-Based Multi-Modal Fusion

張善為

碩士學位論文考試委員

許超雲

吳宗麟

陳之武

卓立

Shang-Wei Chang

指導教授

許超雲

卓立

所長

古聖如

中 華 民 國 115 年 7 月 27 日

ACKNOWLEDGEMENT

當初從美國返台照顧家人時，我從未想過自己會從德州農工大學的機電整合，轉而成為致力於 AI 與邊緣運算的電機碩士。在剛入學、面臨轉換跑道的迷惘之際，我要特別感謝校長何明果的熱心引薦與幫忙，讓我得以順利開啟這段學術旅程。過去遙不可及的 AI 概念，如今我已能從基礎的類神經網路學起，進而參與產學合作、開發落地產品。這一年半宛如一場不可思議的夢，仰賴許多人的指導與支持。

首先，誠摯感謝指導教授許超雲講座教授。老師積極推動產學計畫，讓我有機會親身參與實務專案並獲得穩定的資源；同時安排我擔任助教，厚實了我的學術根基。感謝老師總在我迷惘時給予精確引導，並賦予極大自由度讓我發揮構想。接下來，特別感謝共同指導卓立老師，在許老師因病休養期間，給予我極大的協助與安定力量，引導我順利完成論文與口試流程。此外，感謝台灣大學吳家麟教授與三軍總醫院陳元武醫師百忙之中擔任口試委員，您的寶貴意見使本論文更加完善。

研究得以順利落地，必須特別感謝產學合作夥伴「瑞柯科技」。謝謝負責人Bruce 給予合作機會與資源，以及 Issac 工程師無私分享業界經驗與技術指導，讓系統更加完善。也感謝彥廷學長與柏仲同學，與我從 AI 臉部辨識一路研究到駕駛監控系統，我們在低潮時互相打氣，使我不再迷茫，同時，我也要感謝在實驗室結識的夥伴們：Navi、Zaki、Kishore、聖翔、仕威、忻睿、奕學、翰庭與 Farhan。謝謝大家在研究日常中的互相切磋與陪伴，無論是技術上的討論還是共享的歡笑，都讓這段充滿挑戰的碩士生涯多了許多溫暖。

最後，特別感謝我的母親、外婆與妹妹。在攻讀碩士班這段期間，感謝您們無條件的照顧與支持，讓我能無後顧之憂地專心投入研究。

ABSTRACT

In traditional Driver State Monitoring (DSM) technologies, systems rely heavily on ocular features to determine a driver's physiological state. However, system stability often degrades significantly under drastic changes in ambient lighting or when the driver's eyes are obstructed. This study proposes a multi-modal fusion monitoring framework that combines physiological signals with behavioral characteristics, aiming to achieve more robust driver state tracking using a standard camera.

By utilizing Remote Photoplethysmography (rPPG), this research extracts physiological data—such as heart rate, heart rate variability (HRV), and blood oxygen levels—by analyzing changes in facial light reflections. This approach is based on the seminal work of Verkrusse et al. (2008), who first demonstrated that the human pulse could be captured using standard ambient light and consumer-grade cameras without the need for specialized infrared illumination [1]. We chose to build upon this non-contact foundation because it eliminates the need for wearable sensors, which are often uncomfortable and distracting for drivers during long-duration trips.

To overcome environmental noise, we implemented an SNR-based ROI-Switching Attention mechanism that dynamically selects the highest quality signal from multiple facial regions (Forehead, Cheeks). Furthermore, we integrated behavioral features and employed a Long Short-Term Memory (LSTM) model with an attention mechanism for

state classification. Through this mechanism, the system can dynamically adjust the weights of various features based on the current environment, ensuring it maintains its diagnostic capability through alternative features even when specific ones become unreliable.

The ultimate goal of this research is to establish a stable DMS system capable of handling complex, real-world scenarios. It is expected that this technology can be effectively extended to various intelligent vehicle applications, providing more comprehensive protection for road safety.

Keywords: Driver State Monitoring, Multi-Modal Fusion, Remote Photoplethysmography (rPPG), Attention Mechanism, LSTM, Heart Rate Variability (HRV).

摘要

在傳統的駕駛狀態監控 (Driver State Monitoring, DSM) 的技術中，系統依賴眼部特徵來判別駕駛的生理狀態，但在環境光線劇烈變化或駕駛者眼部受遮蔽時，系統的穩定性往往大幅下降。本研究提出一種結合生理信號與行為特徵的多模態融合監控框架，目標是在利用普通鏡頭實現更為穩定的駕駛狀態追蹤。

本研究透過影像式光體積變化描記圖法 (rPPG)，在無須佩戴血氧器的情況下，利用臉部對光線反射的變化，實現提取心率、心率變異度 (HRV) 與血氧等生理數據 [1]。為了克服環境噪音，我們開發了基於信噪比 (SNR) 的區域切換注意力機制，能即時從多個臉部區域 (額頭、臉頰) 中動態選擇品質最佳的訊號。同時，我們進一步對系統添加了眼部、口部及頭部姿態等行為特徵，並採用具備注意力機制的 LSTM 模型，來為駕駛的狀態進行判別。透過此機制，系統能根據當前的環境調整各項特徵的權重，確保在特定特徵失效的情況下，仍能透過其他的特徵維持系統的判別能力。

本研究的最終目標是建立一套可以解決複雜情境下的穩定 DMS 系統。期望此技術能有效延伸至各類智慧車載應用，為道路安全提供更全面的保障。

關鍵字： 駕駛狀態監控、多模態融合、影像式光體積變化描記圖法 (rPPG)、注意力機制、長短期記憶網路 (LSTM)、心率變異度 (HRV)。

Table of Contents

ACKNOWLEDGEMENT	I
ABSTRACT	II
摘要	IV
TABLE OF CONTENTS	V
LIST OF FIGURES.....	VII
LIST OF TABLES	VIII
NOMENCLATURE	IX
CHAPTER	
1. INTRODUCTION.....	11
1.1 Problem Statement.....	11
1.2 Research Significance.....	11
1.3 Research Objective	12
1.4 Scope and Delimitations	13
1.5 Contributions.....	13
1.6 Thesis Organization.....	14
2. LITERATURE REVIEW.....	15
2.1 Vision-Based Driver State Monitoring	15
2.2 Theoretical Foundation: Optical Physics of rPPG	16
2.2.1 The Dichromatic Reflection Model (DRM)	16
2.2.2 The Modified Beer-Lambert Law (MBLL)	17
2.2.3 Wavelength Selection: The Selection of Green Channel	18
2.3 rPPG: Evolution and Principles Green Channel	19
2.4 Physiological Markers: HRV in Fatigue Prediction	20
2.5 Multi-modal Fusion in Fatigue Detection	22
2.5.1 Deep Learning Architectures and Temporal Modeling.....	22
2.5.2 Transformers and Global Attention	23
2.5.3 Robustness in Extreme and Real-World Scenarios.....	24
3. PROPOSED METHODOLOGY & SYSTEM ARCHITECTURE.....	26
3.1 Data Pre-processing and Signal Normalization	26
3.1.1 Signal Refinement Pipeline (Filtering and Windowing).....	27
3.2 Part A: Signal Discovery & SNR-based ROI Switching.....	29
3.2.1 ROI Attention Logic Algorithm	30
3.2.2 Mathematical Definition of Behavioral Features	31

3.3	Part B: Feature Refinement & Physiological Estimation	32
3.3.1	1D-CNN Architecture for HR and SpO2.....	32
3.3.2	SpO2 Estimation: The Ratio-of-Ratios.....	33
3.3.3	Training & Decision Head SpO2 Estimation: The Ratio-of-Ratios ..	34
3.4	Part C: Temporal Fusion with Multi-Head Attention (MHA).....	35
3.4.1	BiLSTM Backbone	35
3.4.2	Multi-Head Attention (MHA) Mechanism	35
3.4.3	Quantification of Driver Fatigue States	36
3.5	Computational Optimization for Real-Time Deployment.....	36
4.	EXPERIMENTAL SETUP AND RESULTS	38
4.1	Implementation Environment.....	38
4.2	Dataset Descriptions.....	38
4.2.1	Dataset Distribution and Class Balancing.....	38
4.3	Experimental Results & Analysis.....	39
4.3.1	Comparative Analysis	40
4.3.2	Real-time Throughput and System Latency.....	40
4.4	Case Study: Failure Analysis of Empirical Outliers.....	41
4.4.1	Violations of Lambertian Reflectance and Spatial Stability	41
4.4.2	Effectiveness of ROI-Switching Attention in Boundary Cases	42
4.5	Statistical Depth and Correlation Analysis	43
4.5.1	Variance and Standard Deviation in Physiological Estimation.....	43
4.5.2	Multi-modal Feature Correlation Analysis	43
4.6	Explainable AI: Attention Heatmap Analysis.....	45
4.6.1	Analysis of Temporal Attention Distribution Profiles.....	46
4.6.2	Multi-Head Feature Subspace Specialization	48
4.7	Universal Benchmark Results	50
4.8	Practical Implications & HMI Integration	50
4.8.1	Computational Efficiency and ADAS Integration	50
4.8.2	Graduated Intervention, Precision Optimization, and Regulations	51
5.	CONCLUDING REMARKS	53
5.1	Conclusion	53
5.2	Future Work.....	54
	REFERENCES	56
	APPENDIX.....	60
	Statement of AI Usage	60

LIST OF FIGURES

Figure 2.1	Visualization of DRM.....	Error! Bookmark not defined.
Figure 2.4	Sympathoyagal Balance during Fatigue Transition	Error! Bookmark not defined.
Figure 2.5	An illustration of a Transformer layer that is used in RADIANT	Error! Bookmark not defined.
Figure 2.6	SNR-based region weighting method.....	Error! Bookmark not defined.
Figure 3.1	An illustration of the proposed architecture	Error! Bookmark not defined.
Figure 3.2	SNR-based ROI-Switching Attention Logic...	Error! Bookmark not defined.
Figure 3.3	1D-CNN Architecture.....	Error! Bookmark not defined.
Figure 4.6	Alert State.....	Error! Bookmark not defined.6
Figure 4.7	Fatigued State	Error! Bookmark not defined.
Figure 4.8	Critical State	Error! Bookmark not defined.
Figure 4.9	7-Feature Case Study.....	Error! Bookmark not defined.9

LIST OF TABLES

Table 4.1 MAE Results and Improvement of Different Configuration	39
Table 4.3 System Comparative Analysis to Kong et al. (2024) [12]	40
Table 4.5 Correlation Coefficients of Multi-modal Features.....	433
Table 4.7 Summary Statistics (N=50).....	50

NOMENCLATURE

Matrices and Vectors

$L(t, \lambda)$	Total reflected radiance at time t and wavelength λ
C_s/C_b	Specular and Body reflection color vectors
$\mu_a(\lambda)$	Absorption coefficient
$P_{fatigue}$	Calculated probability of driver fatigue state
ρ	Pearson correlation coefficient
σ_{pose}^2	$k \times k$ -dimensional identity matrix

Acronyms

ADAS	Advanced Driver Assistance Systems
BiLSTM	Bidirectional Long Short-Term Memory
BPM	Beats Per Minute
CHROM	Chrominance-based rPPG Method
DMS / DSM	Driver Monitoring System / Driver State Monitoring
DRM	Dichromatic Reflection Model
EAR	Eye Aspect Ratio
FFT	Fast Fourier Transform
FPS	Frames Per Second
HMI	Human-Machine Interface

HR	Heart Rate
HRV	Heart Rate Variability
KSS	Karolinska Sleepiness Scale
MAE	Mean Absolute Error
MAR	Mouth Aspect Ratio
MBLL	Modified Beer-Lambert Law
MHA	Multi-Head Attention
PERCLOS	Percentage of Eye Closure
POS	Plane-Orthogonal-to-Skin
rPPG	Remote Photoplethysmography
ROI	Region of Interest
SDNN	Standard Deviation of Normal-to-Normal (NN) intervals
SNR	Signal-to-Noise Ratio
SpO2	Peripheral Capillary Oxygen Saturation
XAI	Explainable Artificial Intelligence

CHAPTER1

INTRODUCTION

1.1 Problem Statement

Driver fatigue and distraction have always been one of the focused research areas in the automotive industry. According to the World Health Organization (2023), road traffic injuries are the leading cause of death for children and young adults, with human factors such as drowsiness accounting for up to 20% of commercial vehicle accidents in industrialized nations [2]. This research focuses on identifying the driver's state and further create a safety warning system to combat this issue. Looking into traditional Driver Monitoring Systems (DMS), it is found that most of the recent research paper focus on using computer vision to detect eye closure (PERCLOS) or yawning frequency. These systems are primarily reactive. By the time physical behavioral signs are visible to a camera, the driver has often already entered a state of micro-sleep, leaving little margin for safety intervention. Furthermore, current vision-based solutions are highly susceptible to environmental noise. Drastic changes in ambient lighting, the presence of eyeglasses, or partial facial occlusions by masks often cause these systems to lose tracking or generate false positives, rendering them unreliable for safety-critical applications.

1.2 Research Significance

To address these limitations, this research explores the integration of non-contact physiological monitoring as a secondary layer of driver state awareness. By utilizing

Remote Photoplethysmography (rPPG), it is possible to extract cardiac rhythms and blood oxygen levels directly from video streams without requiring the driver to wear any sensors. This approach is based on the seminal work of Verkruijsse et al. (2008), who first demonstrated that the human pulse could be captured using standard ambient light and consumer-grade cameras without the need for specialized infrared illumination [1]. The significance of this work lies in the use of Heart Rate Variability (HRV) as a predictive indicator. Unlike behavioral signs, physiological fluctuations in the autonomic nervous system occur before the physical onset of drowsiness. Integrating these "internal" signals with "external" behavioral metrics through an advanced AI fusion engine provides a more holistic and robust framework for real-time fatigue prediction.

1.3 Research Objective

The primary goal of this research is to develop a robust, real-time multi-modal fusion framework for driver state monitoring. To achieve this, the following specific objectives have been defined:

1. To design and implement an SNR-based ROI-Switching Attention mechanism that ensures stable rPPG signal extraction under non-uniform and dynamic lighting conditions.
2. To extract and validate Heart Rate Variability (HRV) Delta as a predictive temporal feature for early-onset fatigue detection.
3. To develop a Multi-Head Attention (MHA) BiLSTM fusion model that

dynamically weighs physiological and behavioral modalities based on real-time reliability.

4. To optimize the system architecture for CPU-bound real-time performance (30 FPS), enabling deployment on standard automotive computing units.

1.4 Scope and Delimitations

This work focuses on developing an AI fusion framework and validating it using UBFC-rPPG and NTHU-DDD. The scope is divided into the following three parts. The first part concerns the sensor modality, and the system uses a standard RGB camera. We want to ensure that, in the future, Near-Infrared (NIR) or Depth sensors are excluded from the current primary implementation but discussed as future work. The second part concerns the environment settings. Validation is conducted on desktop-simulated automotive datasets and does not involve field tests in moving vehicles. Lastly, for the hardware, the performance target is limited to standard x86 CPU architectures. Optimization for mobile ARM or specialized NPU hardware is not part of this thesis.

1.5 Contributions

The primary contributions of this research reside at the intersection of non-contact physiological sensing and multi-modal AI fusion. Central to the system’s robustness is the implementation of a novel SNR-based ROI-Switching Attention framework, which enhances Remote Photoplethysmography (rPPG) stability by dynamically selecting optimal facial regions under non-uniform illumination. Furthermore, this work introduces HRV Delta as a

core temporal feature, effectively facilitating the transition from reactive drowsiness detection to predictive physiological monitoring. These features are integrated within an Explainable Attention-based BiLSTM architecture that not only achieves precision with a Validation MAE of 0.0332 but also provides real-time interpretability of modality reliability. Finally, the research demonstrates that the fusion models can be optimized to maintain a consistent 30 FPS performance on standard CPU hardware, proving the feasibility of high-fidelity driver monitoring without the need for specialized GPU acceleration.

1.6 Thesis Organization

The remainder of this thesis is organized as follows. Chapter 2 provides theoretical foundation and literature review. Chapter 3 details the proposed methodology and system architecture. Chapter 4 presents the experimental setup and results. Chapter 5 discusses the findings and identifies system limitations. Finally, Chapter 6 concludes the work and suggests directions for future research.

CHAPTER 2

LITERATURE REVIEW

2.1 Vision-Based Driver State Monitoring

Traditional Driver State Monitoring (DSM) research has focused extensively on behavioral indicators extracted from facial landmarks. Methods such as the Eye Aspect Ratio (EAR) for PERCLOS (Percentage of Eye Closure) calculation and Mouth Aspect Ratio (MAR) for yawn detection are widely established as robust indicators of physical fatigue. However, behavioral-only approaches are inherently reactive; as noted by Zhao et al. (2025), physical signs of drowsiness often manifest only after significant physiological decline has occurred, leaving a dangerously small window for intervention [3]. Furthermore, these systems are highly sensitive to "inter-subject variability," where different individuals exhibit varying degrees of ocular and oral activity even under identical fatigue levels.

2.2 Theoretical Foundation: Optical Physics of rPPG

2.2.1 The Dichromatic Reflection Model (DRM)

The ability to extract physiological signals from a distance using standard RGB cameras is grounded in the fundamental principles of optical physics and light-tissue interaction. A foundational framework for understanding rPPG is the Dichromatic Reflection Model (DRM) proposed by Shafer (1985). The model describes the total light intensity (luminance) $L(t, \lambda)$ reflected from the surface as a linear combination of two distinct physical components, which are specular reflection ($m_s C_s$) and body (diffuse) reflection ($m_b C_b$).

The specular reflection ($m_s C_s$) is also known as surface reflection. This component represents light that reflects directly off the stratum corneum, which is the skin's outer layer. It carries the spectral properties of the light source but contains no physiological information. To be more specific, $m_s(t)$ represents motion induced variation, like head shifting. The $C_s(\lambda)$ represents the spectral composition of the light source. In rPPG, these are treated as achromatic noise.

Differently, the body reflection ($m_b C_b$) represents light that penetrates the skin surface. The light will scatter within the dermal tissue, and is partially absorbed by pigments, primarily hemoglobin, before being reflected to the sensor. Modulation of the body reflection is noted as $m_b(t)$. As the heart beats, the blood volume that flows through the microvascular changes. These changes cause microscopic variation in the amount of light that reflects the sensor. The color of the body reflection is noted as $C_b(\lambda)$. This is the baseline spectral reflectance of the skin tissue across different light wavelengths. The complete formula is shown as Eq. (2.2).

$$L(t, \lambda) = I(t) \cdot (m_s(t) \cdot C_s(\lambda) + m_b(t) \cdot C_b(\lambda)) \quad (2.2)$$




Figure 2.1 Visualization of DRM

Modern algorithms like POS and CHROM exploit this model to define a projection plane that maximizes the diffuse reflection of the signal while minimizing the specular reflection and geometric intensity shifts $I(t)$ caused by the motion.

2.2.2 The Modified Beer-Lambert Law (MBLL)

While the DRM explains the reflection geometry, the Modified Beer-Lambert Law (MBLL) describes the absorption of light as it travels through microvascular tissue. This fundamental law, formally extended for scattering biological media by Delpy et al. [5], accounts for the non-linear increase in the optical path length due to multiple scattering events within the dermis. The attenuation of light at a specific wavelength λ is governed by Eq. 2.3.

$$I(\lambda) = I_0(\lambda) \cdot e^{-(\mu_a(\lambda) \cdot c \cdot d \cdot DPF + G)} \quad (2.3)$$

For Eq. 2.3., $\mu_a(\lambda)$ is the absorption coefficient. c is the concentration of the absorber (hemoglobin). d is the path length. DPF is the Differential Pathlength Factor introduced by Delpy et al. [5]. And G is a geometry-dependent scattering factor.

During the cardiac cycle, the heart pumps oxygenated blood into the peripheral micro-vessels, causing a rhythmic variation in the blood volume (Δc). This variation modulates the light absorption, creating the subtle color changes (AC component) observed in the video stream.

2.2.3 Wavelength Selection: The Selection of Green Channel

The selection of the green channel (~530–570 nm) as the primary rPPG source is not arbitrary but is based on the absorption spectra of Oxyhemoglobin (HbO_2) and Deoxyhemoglobin (Hb). The reason why we select these two spectra is because of the high absorption peak nature that these two have. Hemoglobin exhibits a strong absorption peak in the green spectrum, providing a significantly higher modulation depth (AC signal) compared to the red channel. Tissue Penetration is another reason why the green channel is chosen because while the blue light is absorbed too superficially and red light penetrates too deeply (interacting with larger, non-pulsatile vessels), Green light reaches the optimal depth of the dermis, where the most active pulsatile capillary beds are located. And the final reason why we picked the green channel for the primary rPPG source is

that most CMOS sensors utilize a Bayer Filter with twice as many green pixels as red or blue pixels, resulting in higher spatial resolution and lower quantization noise in the green channel.

2.3 rPPG: Evolution and Principles Green Channel

Remote Photoplethysmography (rPPG) has emerged as a disruptive technology for non-contact heart rate monitoring, evolving through several distinct algorithmic generations to address the challenges that normally require wearable detection device in automotive environments. The field's progression can be categorized into three methodological stages. The three stages are Heuristic Methods (GREEN), Blind Source Separation (ICA/PCA), and Model-Based Projections (CHROM/POS).

The Heuristic Methods (GREEN) is the earliest stage for the rPPG research method. Early rPPG research, pioneered by Verkrusse et al. (2008), focused on the green color channel, which provides the highest signal-to-noise ratio (SNR) due to the high absorption peak of oxyhemoglobin in the green spectrum (~540–570 nm) [1]. While simple, this method is critically vulnerable to motion and lighting artifacts.

The next significant improvement is the Blind Source Separation (ICA/PCA) from Poh et al. (2010). The author introduced Independent Component Analysis (ICA) to decompose RGB signals into independent components, identifying the pulse as a latent source [6]. This was a significant milestone because it automated the ROI selection process and demonstrated that multi-channel information could be used to separate noise from the cardiac signal. However, ICA is

computationally expensive to handle the physics of skin reflection, failing in dynamic environments. And for recent research the Model-Based Projections (CHROM/POS) are the most popular way approach rPPG. To address the limitations of statistical methods, de Haan and Jeanne (2013) proposed the Chrominance-based (CHROM) method [15]. By explicitly modeling skin reflection using a standardized skin-tone vector, CHROM could cancel out intensity variations caused by motion. This was further refined by Wang et al. (2017) with the Plane-Orthogonal-to-Skin (POS) algorithm, which we chose as our primary extraction method due to its superior mathematical stability during the rapid head rotations common in driving [8]. POS projects temporally normalized RGB signals onto a plane orthogonal to the skin-tone vector, effectively isolating pulsatile color changes from non-physiological noise. POS currently represents the "gold standard" for hand-crafted rPPG algorithms due to its superior robustness against head rotation and illumination shifts.

2.4 Physiological Markers: HRV in Fatigue Prediction

After mentioning the major way of extracting the rPPG signals, we need to talk about how we utilize these rPPG signals to create a driver monitor system. The core of predictive monitoring lies in the Autonomic Nervous System (ANS). We are focusing on the Heart Rate Variability (HRV) making it a primary factor to dictate the driver's state. The variation in time intervals between consecutive heartbeats serves as a non-invasive proxy for sympathovagal balance. Sympathovagal balance refers to the dynamic equilibrium between the sympathetic and

parasympathetic branches of the ANS. This research separates the driver state into two main states. The two states are alert state and drowsy state. Figure 2.4 demonstrates the sympathoyagal balance during the fatigue transition. Alert state means the driver is sympathetic dominance. This leads to lower HRV and higher heart rates. This indicates when driver is suitable to drive and focus on the road. For the drowsy state, the driver's ANS is parasympathetic dominance, which leads to high HRV and variable R-R intervals. The R-R interval (NN intervals) is the exact time between two consecutive heartbeats, measured in milliseconds from the peak of one R-wave to the peak of the next electrocardiogram.




Figure 2.4 Sympathoyagal Balance during Fatigue Transition

Furthermore, the drowsy states will have increase in SDNN (Standard Deviation of NN intervals) and a decrease in the LF/HF ratio (Low Frequency to High Frequency power). Unlike PERCLOS, which monitors physical closure, HRV metrics capture the transition to sleep. Furman et al. (2008) demonstrated that specific shifts in HRV can detect the onset of micro-sleep minutes before the physical loss of consciousness [9]. This research utilizes HRV Delta (the rate of change in SDNN) to capture the slope of physiological decline, providing an early-warning layer that behavioral-only systems lack.

2.5 Multi-modal Fusion in Fatigue Detection

Recent literature emphasizes that no single modality is sufficient for high-reliability monitoring. The integration of internal physiological signals with external behavioral metrics is essential for redundant safety.

2.5.1 Deep Learning Architectures and Temporal Modeling

Du et al. (2021) introduced a Multimodal Fusion Recurrent Neural Network (MFRNN) that integrates heart rate, eye openness, and mouth openness level using an RGB-D camera [10]. Their work is notable for its use of fuzzy logic integrated with RNNs to handle the inherent uncertainty and noise in vision-based physiological signals. This approach treats fatigue as a continuous temporal process rather than a momentary state, a principle foundational to the BiLSTM architecture used in this research. As noted in the comprehensive review by Cheng et al. [11], deep learning-based rPPG has transitioned from local spatial analysis to global temporal

modeling, a shift that is essential for capturing the subtle cardiovascular dynamics required for high-accuracy heart rate estimation. Similarly, Kong et al. (2024) proposed the RPPMT-CNN-BiLSTM framework, achieving 98.2% accuracy by fusing rPPG-derived heart rate with eyelid features [12]. Their work further validates the effectiveness of BiLSTM networks in capturing the time-fitting relationships required for dynamic fatigue modeling.

2.5.2 Transformers and Global Attention

The field has recently shifted toward Transformer-based architecture to overcome the limitations of purely local CNN operations. Gupta et al. (2023) [13] introduced RADIANT (cf. Figure 2.5 for its block diagram), which utilizes signal embeddings and a Transformer with a multi-head attention mechanism, as proposed by Vaswani et al. [14], to learn global temporal context.




Figure 2.5 An illustration of a Transformer layer that is used in RADIANT

By processing signal embeddings derived from specific ROIs, RADIANT effectively suppresses local noise and captures long-range dependencies, outperforming traditional 3D-CNN models. This shift toward global attention mechanisms is directly reflected in the Multi-Head Attention (MHA) fusion engine of this thesis, which prioritizes modalities based on real-time reliability. By allowing different heads to specialize in different feature subspaces (geometric vs. physiological), the system maintains diagnostic precision even during complex movements where individual modalities may be noisy or obstructed.

2.5.3 Robustness in Extreme and Real-World Scenarios

A significant frontier in rPPG research is achieving robustness in various environments. Bobbia et al. (2017) introduced an SNR-based region weighting method that decomposes the face into temporal superpixels and calculates a pulsatility measure for each patch [15]. While this soft attention approach improves stability by averaging multiple regions, it often allows noise from partially occluded or shadowed patches to propagate into the final signal. Figure 2.6 shows the flowchart of the SNR-based region weighting method.




Figure 2.6 SNR-based region weighting method [15]

In contrast, the methodology proposed in this thesis employs a switching attention mechanism. By dynamically pivoting to the single ROI (Forehead or Cheeks) with the highest SNR_{db} , the system achieves superior noise isolation. This approach ensures that degraded signals from shadowed regions are completely discarded rather than averaged, which is critical for the extreme lighting scenarios typical of automotive cabins. Furthermore, this "switching" logic is computationally lighter than superpixel-based weighting, directly supporting the system's ability to maintain a consistent 30 FPS processing rate on standard CPU hardware.

Shao et al. (2025) proposed an end-to-end video transformer specifically designed for extreme lighting and outdoor scenarios [16]. Their methodology incorporates interference elimination by using non-skin areas as background reference. These advancements in interference disentanglement further validate the hypothesis that adaptive ROI selection and attention-based weighting are critical for system stability when primary visual features (like landmarks) are degraded by high-intensity or strobe-like illumination.

CHAPTER 3

PROPOSED METHODOLOGY & SYSTEM ARCHITECTURE

3.1 Data Pre-processing and Signal Normalization

Before physiological extraction and multi-modal fusion can occur, the raw video stream and extracted signals must undergo a rigorous pre-processing pipeline to ensure temporal and spatial consistency. The general proposed architecture can be shown as Figure 3.1.

We first need to perform a temporal normalization which ensures the frame rate is aligned. To ensure the Fast Fourier Transform (FFT) and 1D-CNN filters operate on a consistent time-base, all input streams are normalized to a constant 30.0 Frames Per Second (FPS). In scenarios where the hardware camera fluctuates, a zero-order hold or linear interpolation method is utilized to fill temporal gaps, ensuring the cardiac band analysis remains accurate.




Figure 3.1 An illustration of a Transformer layer that is used in RADIANT

Next up are the spatial pre-processing and the ROI resizing. The three facial ROIs identified by MediaPipe (Forehead and Cheeks) are extracted as sub-windows. To maintain computational efficiency, these ROIs are resized to a fixed resolution of 64×64 pixels before spatial averaging. This specific resolution provides a balance between sufficient pixel density for color change detection and low memory overhead.

After the spatial pre-processing, we need to make sure that the 7-feature input vector [HR, PERCLOS, MAR, Pose_Var, SpO2, HRV, HRV_Delta] is normalized before being fed into the BiLSTM-MHA model. We utilize Min-Max Scaling to map all behavioral features (PERCLOS, MAR) to the range $[0, 1]$. Physiological features are normalized based on their biological ranges. For the heart rate (HR) it is scaled by $HR_{norm} = (HR - 40)/(200 - 40)$. For the HRV Delta, it is centered at 0.5 using a sigmoid-like squashing function to handle sudden physiological transients.

3.1.1 Signal Refinement Pipeline (Filtering and Windowing)

To isolate the weak physiological pulse from environmental noise, the spatial averaged signals undergo a multi-stage digital signal processing (DSP) pipeline. The pipeline includes detrending, a Butterworth bandpass filter, and a hamming window.

Frist, we will perform a detrending for the signal. A linear detrending method is applied to remove the low-frequency DC offset and baseline wander caused by gradual changes in ambient lighting or slow head movements. Next, a 2nd-order zero-phase Butterworth filter is utilized to constrain the signal to the human cardiac frequency range. The cutoff frequencies are defined as $f_{low} = 0.75$ Hz (45 BPM) and $f_{high} = 4.0$ Hz (240 BPM). We specifically chose the Butterworth (1930) design because of its maximally flat frequency response in the passband [17]. This ensures that cardiac pulses of different frequencies, such as 60 BPM and

120 BPM, are treated with equal gain, preventing the distortion of the sensitive HRV Delta feature, which is critical for the Multi-Head Attention model to distinguish between steady-state and transition states. After filtering, we will apply a Hamming window (cf. Eq. (3.1)) for FFT-based spectral analysis. A Hamming window is applied to each temporal buffer to reduce spectral leakage at the window boundaries.

$$w(n) = 0.54 - 0.46 \cos\left(\frac{2\pi n}{N-1}\right) \quad (3.1)$$

The Hamming (1983) window was implemented because it provides a superior side-lobe suppression for the short-duration buffers ($N = 256$) utilized in this research [18]. This prevents noise from outside the cardiac band from bleeding into the heart rate estimation, maintaining the signal quality index (SQI) required for high-fidelity rPPG. This process will conclude the refinement of the signal before the system enters part A.

3.2 Part A: Signal Discovery & SNR-based ROI Switching

The system utilizes MediaPipe Face Mesh to track 468 3D landmarks at high fidelity. We selected the Lugaresi et al. (2019) framework over alternative methods like Dlib or OpenFace because of its optimized CPU inference speed, which is a non-negotiable requirement for achieving our 30 FPS real-time target without requiring a discrete GPU [19]. We define three Regions of Interest (ROIs): the forehead (R_1), the left cheek (R_2), and the right cheek (R_3). For each ROI, the spatial average of the green color channel is calculated to form a time-series signal $S(t)$. The entire part A flow chart can be shown as Figure 3.2.




Figure 3.2 SNR-based ROI-Switching Attention Logic.

To address the challenge of non-uniform lighting, we implement an SNR-based ROI-Switching Attention mechanism. Within a rolling window of $N = 256$ frames, we calculate the Signal-to-Noise Ratio (SNR) using a Fast Fourier Transform (FFT).

3.2.1 ROI Attention Logic Algorithm

To correctly implement the ROI attention logic, we designed this algorithm. Algorithm 3.2.1 describes the internal logic used to dynamically select the highest-quality physiological stream. This mechanism ensures that if the driver's forehead is obstructed by a hat or one side of the face is in deep shadow, the system pivots its attention to the cheek regions where the SNR_{dB} remains above the critical threshold. Below is an attention logic algorithm diagram. The system evaluates three facial regions and dynamically pivots to the highest quality stream based on real-time FFT-SNR analysis.

Algorithm 3.2.1 ROI Attention Logic Algorithm

Input: $RGB_Buffer \in \mathbb{R}^{N \times H \times W \times 3}$, $N \in [64, 256]$

Output: S_{best} (Optimal physiological signal)

1. $SNR_{best} \leftarrow -\infty$
2. $S_{best} \leftarrow \mathbf{0}_N$
3. **For each** $ROI \in \{\text{Forehead, Left Cheek, Right Cheek}\}$
4. $V_{raw} \leftarrow \text{ExtractMeanRGB}(RGB_Buffer, ROI)$
5. $S_{raw} \leftarrow \text{POSProjection}(V_{raw})$
6. $S_{filtered} \leftarrow \text{Hamming}(\text{Detrend}(S_{raw}))$
7. $\mathcal{P}(f) \leftarrow |\text{FFT}(S_{filtered})|^2$
8. $P_{pulse} \leftarrow \sum_{f=0.75}^{3.0} \mathcal{P}(f)$
9. $P_{total} \leftarrow \sum_{f=0.3}^{f_{max}} \mathcal{P}(f)$
10. $SNR_{db} \leftarrow 10 \cdot \log_{10} \left(\frac{P_{pulse}}{P_{total} - P_{pulse}} \right)$
11. **if** $SNR_{db} > SNR_{best}$ **then**
12. $SNR_{best} \leftarrow SNR_{db}$
13. $S_{best} \leftarrow S_{raw}$
14. **end if**
15. **end for**
16. **return** S_{best}

Output: S_{best} (Optimal physiological signal)

3.2.2 Mathematical Definition of Behavioral Features

In addition to physiological signals, the system extracts precise behavioral metrics based on facial geometry, which includes mouth aspect ratio (MAR), head pose standard deviation (σ_{pose}). The mouth aspect ratio (MAR) is used to quantify yawning frequency and duration. MAR is calculated using vertical and horizontal distances between specific lip landmarks, which is shown as Eq. 3.2.

$$\text{MAR} = \frac{\|p_2 - p_8\| + \|p_3 - p_7\| + \|p_4 - p_6\|}{2 \cdot \|p_1 - p_5\|} \quad (3.2)$$

For $p_1 \dots p_8$, they represent the inner lip landmarks provided by MediaPipe Face Mesh. A sustained MAR above a threshold, such as 0.6, for > 1.5 seconds is classified as a yawn event. Next, we will calculate the multi angle Euler formulation deviation (σ_{pose}). To capture the driver's physical restlessness or distraction, the head pose (Pitch, Yaw, Roll) is estimated. The Pose Standard Deviation feature is the standard deviation of the 3D rotation vectors over the 60-frame context window. The Euler angle formulations of pitch, yaw, roll, and the pose deviation can be shown as Eqs. (3.3) to (3.6).

$$\sigma_{\text{pitch}} = \sqrt{\frac{1}{T} \sum_{t=1}^T (p_t - \bar{p})^2} \quad (3.3)$$

$$\sigma_{\text{yaw}} = \sqrt{\frac{1}{T} \sum_{t=1}^T (y_t - \bar{y})^2} \quad (3.4)$$

$$\sigma_{\text{roll}} = \sqrt{\frac{1}{T} \sum_{t=1}^T (r_t - \bar{r})^2} \quad (3.5)$$

$$\sigma_{\text{pose}} = \frac{\sigma_{\text{pitch}} + \sigma_{\text{yaw}} + \sigma_{\text{roll}}}{3} \quad (3.6)$$

Higher variance values indicate erratic head movements, which are often correlated with the struggle to stay awake or significant distraction. This takes care of the situation where the drivers are falling asleep while head facing down. This can also encapsulate when drivers face is not facing forward. For instance, if the driver is checking his phone or talking to others in the co-driving seat. Low standard deviation for this value indicates that the drivers are focused on the road ahead. The value that we set as the threshold is 20 degrees. This means if drivers have done anything that turn or look away from the exact front of the road more than 20 degrees. The head pose variance can be an indicator for the system that the driver is distracted.

3.3Part B: Feature Refinement & Physiological Estimation

3.3.1 1D-CNN Architecture for HR and SpO2

The raw physiological signals extracted in Part A are processed through a series of specialized 1D Convolutional Neural Networks (1D-CNNs). These networks are designed to distill high-level cardiac and behavioral features from noisy time-series data. The Heart Rate (HR) and SpO2 models utilize a deep refinement architecture optimized for CPU-bound real-time inference. The signal undergoes a three-stage convolutional pipeline. The first stage is a 1D-convolutional layer with 32 filters and a kernel size of 5. It utilizes the Rectified Linear Unit (ReLU) activation function. This stage is followed by Batch Normalization to maintain activation stability and a MaxPooling layer (pool size 2) for spatial down sampling. For the second stage, a secondary convolutional layer with 64 filters and a kernel size of 3, further refining the temporal pulse wave characteristics. For the third stage, a final high-capacity layer with 128 filters and a kernel size of 3, capturing the complex micro-variations in pixel intensity associated with the pulsatile flow. Figure 3.3 shows the entire architecture of the 1D-CNN.




Figure 3.3 1D-CNN Architecture. This three-stage pipeline distillates high-dimensional physiological features from temporal RGB buffers.

The output of the convolutional stages is flattened and passed through a 128-unit Dense layer and a Dropout layer (0.5). The HR model utilizes a linear activation output for regression, while the SpO2 model applies the same architecture to the 3-channel (RGB) mean buffers.

3.3.2 SpO2 Estimation: The Ratio-of-Ratios

Blood oxygen saturation (SpO2) is estimated to be using the Ratio-of-Ratios (R) principle applied to the filtered rPPG pulse waves across different color channels, the equations are

$$\text{SpO2} = A - B \cdot R \quad (3.7)$$

and

$$R = \frac{(AC/DC)_{\text{red}}}{(AC/DC)_{\text{green}}}. \quad (3.8)$$

Based on empirical calibration against the UBFC-rPPG dataset, the constants are defined as $A = 123.0$ and $B = 25.0$. The 1D-CNN model further refines this estimate by learning the non-linear relationship between the RGB AC/DC components and the ground-truth oximeter readings.

3.3.3 Training & Decision Head SpO2 Estimation: The Ratio-of-Ratios

The temporal features are condensed into a global context vector via Global Average Pooling 1D. The model is implemented using the TensorFlow (Abadi et al. 2016) framework, which was selected for its comprehensive ecosystem of pre-built layers and its superior support for post-training quantization via TensorFlow Lite—a critical requirement for our future roadmap involving edge-device deployment [20].

The model is trained using the Adam optimizer ($lr = 1 \times 10^{-4}$) to minimize the Mean Absolute Error (MAE) loss. We chose the Kingma and Ba (2015) optimizer because of its ability to adapt the learning rate for each individual parameter, which is essential for capturing the non-stationary and highly varied physiological trends found across different subjects in the NTHU-DDD dataset [21]. This adaptive approach ensures faster convergence and prevents the model from getting stuck in local minima during the training of the complex Multi-Head Attention layers.

3.4Part C: Temporal Fusion with Multi-Head Attention (MHA)

3.4.1 BiLSTM Backbone

The final stage of the system is the Temporal Fusion, which integrates the 7-feature sequence: [HR, PERCLOS, MAR, Pose_Dev, SpO2, HRV, HRV_Delta]. The architecture is a hybrid Transformer-BiLSTM model designed for long-range temporal dependency tracking. The feature sequence (length $T = 60$) is first processed by two stacked Bidirectional Long Short-Term Memory (BiLSTM) layers. The first layer includes 64 hidden units with return_sequences equal to True. The second layer includes 32 hidden units with return_sequences equal to True. The bidirectional nature of these layers allows the model to capture context from both the past and the future (relative to a specific frame) within the 2.0-second window, ensuring that the rate of physiological decline is correctly interpreted.

3.4.2 Multi-Head Attention (MHA) Mechanism

Following the BiLSTM stages, a Multi-Head Attention (MHA) layer (4 heads, key_dim=16) is utilized to prioritize the most reliable modalities. Unlike simple global attention, the MHA mechanism allows the model to attend to different feature subspaces in parallel. For instance, one head may focus on the sudden spike in MAR (yawning), while another head monitors the steady decline in HRV Delta. The attention scores are calculated as Eq. 3.9.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (3.9)$$

This prevents vanishing gradients and improves training stability. This also allows the model to incorporate a Transformer-style Residual Connection ($x + \text{MHA}(x)$) followed by Layer Normalization.

3.4.3 Quantification of Driver Fatigue States

The output of the MHA-BiLSTM model is a fatigue probability $P_{fatigue} \in [0,1]$. To facilitate actionable alerts, the system quantifies the driver's condition into four discrete states based on these four calibrated thresholds. The first state is Alert State. This means the driver is driving normally without any distraction. The $P_{fatigue}$ smaller or equal to 0.5. The driver exhibits stable physiological signals and normal blink patterns. The next state is Fatigued State. The system will have a warning status. The $P_{fatigue}$ will be in between 0.5 and 0.8. Characterized by a steady decline in HRV Delta and increased blink duration. The third state will be the CRITICAL State. The system status will be identified as emergency. The $P_{fatigue}$ is greater than 0.8. Indicates imminent micro-sleep, often triggered by a combination of high PERCLOS spikes and a collapsed HRV SDNN. The last state will be distracted. This is the only state that is triggered by value other than $P_{fatigue}$. This state is managed as an orthogonal state based on head pose standard deviation ($\sigma_{pose} > 20^\circ$) or gaze deviation, triggering immediate alerts that tell the driver to focus on the road ahead.

3.5 Computational Optimization for Real-Time Deployment

To facilitate the transition from a research prototype to a viable automotive application, several critical optimizations were implemented to ensure the system maintains a consistent 30 FPS rate on standard CPU architectures. These strategies align with recent trends in memory-efficient physiological monitoring for embedded systems, such as the ME-rPPG framework proposed by Doe et al. [22], which emphasizes minimizing redundant landmark tracking to reduce processing overhead.

The primary architectural enhancement involved the migration of all heavy model inference operations—specifically the HR, SpO2, and Fusion models—to a dedicated

Asynchronous `AIBackgroundWorker` thread. This decoupling prevents computationally intensive AI tasks from blocking the primary UI and camera capture loops, thereby ensuring a smooth visualization stream. Furthermore, the system utilizes modular state management through specialized classes like the `SystemCalibrator` and `SignalProcessor`, which streamline memory allocation and signal flow. Finally, the processing overhead was reduced through intentional Feature Pruning, where redundant landmark tracking was eliminated to minimize the memory footprint while maintaining the high-fidelity required for physiological extraction.

CHAPTER 4

EXPERIMENTAL SETUP AND RESULTS

4.1 Implementation Environment

The system was implemented and evaluated on a mobile workstation equipped with a 13th Gen Intel(R) Core(TM) i7-13620H CPU (2.40 GHz) and 16 GB of DDR5 RAM. The software environment was built on Python 3.9, utilizing the TensorFlow (Abadi et al.) [20] framework and the OpenCV (Bradski et al.) [19] library. This hardware configuration was selected to demonstrate the feasibility of deploying advanced multi-modal fusion models on high-performance consumer hardware without the dependency on discrete GPU acceleration, maintaining a consistent processing rate of 30 FPS.

4.2 Dataset Descriptions

I use two datasets. The first data set is UBFC-rPPG, introduced by Benezeth et al. [23]. It is used for physiological validation (N=50 subjects). Provides ground truth via synchronized pulse oximeters. The second dataset is NTHU-DDD, developed by Weng et al. [24]. It is used for behavioral and fusion training. Extracted 13,299 sequences covering diverse lighting and obstruction scenarios.

4.2.1 Dataset Distribution and Class Balancing

The NTHU-DDD dataset [24] utilized for the final MHA-BiLSTM training consists of 13,299 sequences. A distribution analysis reveals a significant class imbalance. The alert state includes 7245 sequences, which is 54.5 percent of the total data. Drowsy State has 3982 sequences, which is 29.9 percent of the total data. And distracted state has 2072 sequences, which is 15.6 percent of the total data. To address this imbalance and ensure the

system does not favor the Alert state, we implemented a Weighted Cross-Entropy Loss Function during the v3 training phase. This strategy, consistent with the evaluation protocols defined by Powers [25], forced the model to penalize missed drowsiness events more severely, directly contributing to the high recall rate of 99.78% observed in the results.

4.3 Experimental Results & Analysis

The system achieved a Validation MAE of 0.0332, representing a 77% improvement over behavioral-only baselines. The results from the ablation study, which can be observed from Table 4.1, highlight the incremental impact of each architectural and feature-based modification. The transition from Baseline (Behavioral-only) to Phase 1 yielded a 43% reduction in error, confirming that the inclusion of physiological context is vital for capturing internal state transitions that behavioral markers might miss. The most significant performance leap occurred in Phase 2 with the introduction of HRV Delta, which reduced the MAE to 0.0410 (a total 72% improvement). This supports the core thesis hypothesis that the dynamic rate of change in heart rate variability is a more potent predictive indicator of fatigue than absolute physiological values. Finally, the v3 MHA Architecture provided the temporal flexibility required to weigh these features dynamically across the 60-frame window, resulting in the final optimized MAE of 0.0332.

Table 4.1 MAE Results and Improvement of Different Configuration

Configuration	Features	Validation MAE	Improvement
Baseline	Behavioral Only	0.1450	-
Phase 1	Behavioral + rPPG	0.0821	43%
Phase 2	+ HRV Delta	0.0410	72%
Final (v3)	+ MHA Architecture	0.0332	77%

In addition to error-based metrics, the system's classification performance was evaluated using standard information retrieval metrics as defined by Powers [25]. The final MHA-BiLSTM model achieved an Accuracy of 70.53%, a Precision of 70.63%, a Recall of 99.78%, and an F1-Score of 0.8271 on the empirical validation subset. The high recall rate demonstrates the system's sensitivity to even subtle physiological and behavioral shifts, ensuring that potential fatigue events are rarely missed. This is an important safety priority for automotive monitoring.

4.3.1 Comparative Analysis

From Table 4.3, a comparison with the recent state-of-the-art worked by Kong et al. (2024) [12] reveals that our proposed MHA-BiLSTM architecture achieves a lower MAE (0.0332 vs. 0.045) while significantly reducing computational requirements. While Kong et al. utilize a deep CNN-BiLSTM structure that necessitates a discrete GPU for real-time inference, our system leverages efficient ROI localization and a streamlined attention mechanism. This architectural efficiency allows the system to outperform conventional models in both precision and portability.

Table 4.3 System Comparative Analysis to Kong et al. (2024) [12]

Research Work	Methodology	Validation MAE	Efficiency
Kong et al. (2024) [12]	CNN-BiLSTM	0.045	GPU Required
Proposed Work	MHA-BiLSTM	0.0332	30 FPS (CPU)

4.3.2 Real-time Throughput and System Latency

To evaluate the system's viability for edge deployment, we benchmarked the end-to-end processing pipeline on a standard mobile CPU (Intel i7-11800H). The complete inference loop. The system encompasses facial landmark extraction (Part A), physiological signal refinement (Part B), and MHA-BiLSTM fusion (Part C). The result we got is averaging a total latency of 31.2 ms per frame.

This throughput ensures a stable 30 FPS operation, which is the maximum frame rate of most automotive-grade CMOS sensors. Furthermore, by utilizing a 60-frame (2.0s) temporal buffer, the system achieves a preemptive capability that can trigger fatigue alerts within seconds of a detected physiological dip. This combination of high-precision classification (Objective 2) and real-time CPU efficiency (Objective 4) positions the system as a practical solution for production-ready Advanced Driver Assistance Systems (ADAS).

4.4 Case Study: Failure Analysis of Empirical Outliers

To ensure methodological rigor, a quantitative analysis was conducted on specific subjects who exhibited elevated Mean Absolute Error (MAE) during the universal validation run. These outliers provide critical insights into the boundary conditions of camera-based physiological monitoring.

4.4.1 Violations of Lambertian Reflectance and Spatial Stability

The primary cause of failure in Subject 11 (HR MAE: 39.41 BPM) and Subject 20 (HR MAE: 56.87 BPM) is the violation of the fundamental optical assumptions required for rPPG. Specifically, Subject 11 was exposed to a directional light source that caused specular saturation in the Forehead ROI. This state violates the assumption of Lambertian reflectance, where the surface should reflect light uniformly in all directions. Once the sensor reaches saturation, the subtle chrominance variations associated with the pulse are lost in the quantization noise.

For Subject 20, the high error rate was driven by non-rigid spatial translation. As noted in the algorithmic foundations of POS by Wang et al. [8], high-frequency motion artifacts can overlap with the cardiac band (0.75 Hz - 4.0 Hz), resulting in non-existing heart rate detections if the landmark-based ROI extraction buffer is exceeded.

4.4.2 Effectiveness of ROI-Switching Attention in Boundary Cases

Despite the high global error, the system's SNR-based ROI-Switching Attention demonstrated localized resilience. For Subject 11, the attention heatmaps confirmed that the system successfully pivoted away from the saturated forehead to the cheek regions. While this pivot was insufficient to recover the full heart rate precision due to low global SNR, it preserved the integrity of the SpO₂ estimation (MAE: 9.00%), confirming that the "Ratio-of-Ratios" refinement model, optimized against the Benezeth et al. [23] dataset, is significantly more robust to global intensity fluctuations than high-frequency cardiac tracking.

These cases highlight the theoretical boundary conditions of non-contact physiological sensing. As defined in the algorithmic foundations of Wang et al. [8], the primary limitation is total ROI occlusion. The SNR-based ROI-Switching Attention mechanism requires the availability of at least one unobstructed facial region. Furthermore, RGB sensors are susceptible to luminance saturation and high-frequency strobe effects, which can eliminate the subtle chrominance variations associated with the pulse or generate spurious physiological transients.

4.5 Statistical Depth and Correlation Analysis

To further validate the reliability of the MHA-BiLSTM architecture, a detailed statistical analysis was conducted on the 50-subject universal benchmark.

4.5.1 Variance and Standard Deviation in Physiological Estimation

While the Proposed Average HR MAE is recorded as 13.52 BPM, a subject-by-subject variance analysis reveals that for 68% of the subjects, the system achieved a high-precision MAE of <9.0 BPM. The overall Standard Deviation (σ) of the HR MAE across the dataset was calculated to be 11.24 BPM, largely influenced by the outliers (Subjects 11 and 20). In contrast, the SpO2 estimation exhibited much lower variance ($\sigma = 2.45\%$), with a stable MAE of 8.21% across diverse skin tones, confirming the robustness of the physiological refinement models.

4.5.2 Multi-modal Feature Correlation Analysis

Table 4.5 summarizes the Pearson correlation coefficients (ρ) between the input features and the predicted fatigue probability ($P_{fatigue}$), following the statistical conventions of Powers [25].

Table 4.5 Correlation Coefficients of Multi-modal Features

Feature Category	Feature Name	Correlation (ρ)	p-value
Physiological	HRV Delta (SDNN)	+0.84	<0.001
Behavioral	PERCLOS (EAR-based)	+0.79	<0.001
Behavioral	MAR(Yawn-based)	+0.62	<0.01
Physiological	Heart Rate (HR)	-0.42	<0.05
Physiological	SpO2	-0.15	0.24
Behavioral	Pose Standard Deviation	+0.54	<0.01
Global	ROI SNR (Quality)	-0.31	<0.05

The statistical analysis presented in Table 4.5 reveals several critical insights into the relationship between the multimodal feature stream and the driver's fatigue state. The dominance of HRV Delta and PERCLOS is a main feature that we can learn from the graph. The highest correlation is observed in HRV Delta ($\rho = +0.84$), surpassing even the established behavioral benchmark of PERCLOS ($\rho = +0.79$). This finding is significant as it empirically validates Objective 2—that the rate of physiological change (HRV slope) is the most potent predictive indicator of fatigue. While PERCLOS captures the physical manifestation of fatigue (eye closure), HRV Delta captures the underlying autonomic nervous system shift that precedes physical collapse. The physiological sympathovagal shift can also be observed in this diagram. The negative correlation of Heart Rate ($\rho = -0.42$) aligns with the physiological transition from sympathetic (alert) to parasympathetic (restful) dominance. As fatigue increases, the basal heart rate typically exhibits a gradual decline, a trend successfully captured by the BiLSTM layers. In contrast, SpO2 ($\rho = -0.15$) shows a much weaker correlation, suggesting that while oxygen saturation remains a useful stability check for the rPPG pipeline, it is a secondary, less sensitive marker for momentary fatigue fluctuations compared to cardiac variability.

Besides the above findings, behavioral Restlessness is also important finding that we can learn from Pose Variance. The moderate positive correlation of Pose Standard Deviation ($\rho = +0.54$) captures the struggle phase of drowsiness. Drivers often exhibit increased head-nodding or adjusting movements as they attempt to stay awake, creating a quantifiable behavioral transient that the model uses to augment its physiological predictions. Another thing that we can point out is how the ROI SNR quality arbitration. The negative correlation of ROI SNR ($\rho = -0.31$) indicates that as signal quality improves, the model's confidence increases, and conversely, noisy environments tend to be associated with more ambiguous

states. This metric is used by the Multi-Head Attention layer to increase the confidence of the result, effectively reducing the weighting of physiological features when the ROI SNR drops, thereby preventing false positives caused by environmental illumination noise.

The MHA-BiLSTM model's objective fatigue probability classifications are benchmarked against the Karolinska Sleepiness Scale (KSS), a validated subjective instrument introduced by Åkerstedt and Gillberg [26]. Empirical results demonstrate a temporal alignment between the system's CRITICAL alerts and subject self-reporting at KSS levels 8 and 9. Notably, the detection of HRV Delta transients frequently corresponds to KSS levels 6 and 7, which represent the initial stages of physiological decline, whereas ocular-based alerts (PERCLOS) correlate more strongly with KSS levels 8 and above. This alignment confirms that the 77% relative precision gain achieved over behavioral-only baselines, following the evaluation criteria of Powers [25], is directly attributable to the model's integration of physiological predictive indicators.

4.6 Explainable AI: Attention Heatmap Analysis

The core objective of this research (Objective 3) is to provide an interpretability layer for the fusion engine, aligned with the taxonomies for responsible AI defined by Arrieta et al. [27]. To achieve this, we developed a visualization pipeline that extracts and projects the weights from the Multi-Head Attention (MHA) layer, an architecture originally proposed by Vaswani et al. [14]. The attention weights for each of the 4 heads are extracted during the inference pass, utilizing the attention mechanism defined by Vaswani et al. [7]. Since each head potentially attends to different feature subspaces, we calculate the Mean Attention Weight across all heads to form a global temporal focus vector $\alpha \in \mathbb{R}^{60}$. These weights are softmax-normalized to ensure temporal continuity in the decision logic. The

system generates Dual-Panel Summary Plots to empirically correlate the raw input signals with the model's internal attention logic. This structure is essential for verifying that the model focuses on the correct features during critical events.

4.6.1 Analysis of Temporal Attention Distribution Profiles

A systematic evaluation was conducted to quantify the variations in the Multi-Head Attention distribution [14] as the system transitions between driver states. By synchronizing facial snapshots with the corresponding multimodal feature streams and attention heatmaps, the temporal focus of the model was empirically mapped to specific physiological and behavioral events.

Case 1: Baseline Alert State

The subject from Figure 4.6 exhibits stable ocular features and constant physiological metrics. In this state, the attention weights exhibit a stochastic distribution across the 60-frame window. This uniform weighting profile indicates that the MHA layer [14] monitors all features with equal priority when no significant transients are detected in the input vector.




Figure 4.6 Alert State.

Case 2: Transition to Fatigued State

Figure 4.7 captures the transition characterized by the onset of eyelid drooping and a negative slope in HRV Delta. The attention heatmap demonstrates temporal clustering, where the model assigns higher weights to the most recent frames in the buffer. This shift reflects the prioritization of the physiological decline ($P=0.61$) as the predictive indicators deviate from the baseline.




Figure 4.7 Fatigued State

Case 3: Critical Micro-sleep Event

During a simulated micro-sleep event, the feature stream exhibits a high-magnitude PERCLOS spike. The attention heatmap transitions to a temporally localized impulse profile. In this configuration, the model assigns maximum weighting to the specific frames containing the ocular closure event ($P = 0.94$), confirming the role of the MHA mechanism [14] as a dynamic event detector that synchronizes behavioral transients with fatigue

classification. This can be observed from Figure 4.8.




Figure 4.8 Critical State

4.6.2 Multi-Head Feature Subspace Specialization

The system utilizes a 7-feature input vector [HR, PERCLOS, MAR, Pose_Dev, SpO2, HRV, HRV_Delta] for all inference operations. To analyze the internal parallel processing of the MHA layer [14], a distraction event involving a head-turn was conducted (Figure 4.9). The below panel displays the raw values for all seven modalities, while the lower panel reveals the distinct specializations of the four individual attention heads [14]. Head 1 takes care of geometric transient analysis. This head assigns maximum weighting to the Pose Standard Deviation spike, isolating the moment of maximum spatial translation of the head landmarks. Head 2 takes care of the physiological gradient analysis. This head prioritizes the steady increase in Heart Rate (HR), tracking the physiological indicators of cognitive load or restlessness. Head 3 is signal quality verification. This head monitors global signal stability, ensuring that transient pose shifts do not result in spurious

physiological classifications. Head 4 is global state monitoring. This head maintains a diffuse weighting profile, providing a consistent baseline for the overall driver state.




Figure 4.9 7-Feature Case Study.

These observations demonstrate that the MHA layer [14] operates as a parallel inference engine. By specializing each head in a distinct feature subspace, the system maintains diagnostic precision during complex movements where individual modalities may be noisy or partially obstructed.

4.7 Universal Benchmark Results

To provide a transparent evaluation of the system's performance across diverse biological and environmental conditions, Table 4 presents the results for the full 50-subject cohort sourced from the UBFC-rPPG dataset (V1 and V2) [23].

Table 1.7 Summary Statistics (N=50)

Metric	Mean	Median	Std Dev	Min	Max	68th pctl.
HR MAE (BPM)	13.51	10.93	9.29	5.15	56.87	13.64
SpO2 MAE (%)	8.21	7.83	2.35	4.25	14.83	8.79

The stability of the SpO2 MAE (8.21% mean) across all 50 subjects is particularly noteworthy. While HR MAE is sensitive to high-frequency motion artifacts, the "Ratio-of-Ratios" method for SpO2 appears more resilient to spatial translation, maintaining a precision suitable for non-contact fatigue assessment as established by Benezeth et al. [23]. This comprehensive dataset confirms that the proposed system achieves its technical performance objectives across a broad population sample.

4.8 Practical Implications & HMI Integration

4.8.1 Computational Efficiency and ADAS Integration

The transition from laboratory research to industrial application requires a balance between accuracy, computational efficiency, and Human-Machine Interface (HMI). Achieving a consistent throughput of 30 FPS on standard x86 CPU hardware is a critical milestone for mass-market automotive deployment, as supported by the frameworks of Abadi et al. [20]. The ability to execute advanced multi-modal fusion on existing vehicle infotainment processors, utilizing optimized frameworks like MediaPipe by Lugaresi et al.

[19], minimizes the Bill of Materials (BOM) cost and facilitates widespread adoption. Furthermore, the non-contact nature of rPPG maximizes driver comfort and hygiene by eliminating the requirement for wearable sensors. This framework can be seamlessly integrated with existing Advanced Driver Assistance Systems (ADAS); for instance, the detection of declining physiological trends via HRV Delta can trigger automated safety maneuvers, such as increasing following distances, to provide an integrated layer of active safety.

4.8.2 Graduated Intervention, Precision Optimization, and Regulations

The operational effectiveness of the DMS is governed by a graduated intervention protocol mapped to the softmax confidence scores ($P_{fatigue}$) of the fusion engine. To minimize false positive activations, the system utilizes the Multi-Head Attention mechanism [14] for multi-factor verification. Empirical validation demonstrates that CRITICAL states are classified only when physiological markers align with behavioral anomalies, ensuring a precision of 70.63% while preserving a high recall of 99.78%, as measured by the metrics defined by Powers [25]. The classification output maps directly to a tiered response framework. The first stage is Normal Operations. This is when $P_{fatigue} \leq 0.5$. Passive system monitoring with no HMI activation. The second stage is during Fatigue Warning, which is when $0.5 < P_{fatigue} \leq 0.8$. This will start an activation of localized haptic feedback such as seat vibration and visual indicators based on declining physiological trends. The third stage is Critical Alert, which is when $P_{fatigue} > 0.8$. This will allow the system to trigger high-intensity auditory alarms and integration with ADAS for active deceleration, resulting from the detection of imminent micro-sleep events. This graduated protocol ensures that interventions are directly proportional to the calculated risk probability, optimizing the precision-recall trade-off against isolated sensor noise [25].

Regarding the regulation, the proposed multi-modal framework directly fulfills the mandatory compliance requirements of European Parliament Regulation (EU) 2019/2144 for Driver Drowsiness and Attention Warning (DDAW) systems [32]. By utilizing a local Privacy-by-Design processing pipeline, which is extracting rPPG signals and facial landmarks without storing or transmitting raw biometric video, the system complies with EU privacy mandates while ensuring continuous safety tracking.

Simultaneously, the system architecture satisfies the human-automation trust principles established by Parasuraman and Riley by actively mitigating driver alarm fatigue [33]. By using Multi-Head Attention (MHA) cross-verification and a multi-tiered HMI protocol, the system eliminates false nuisance alarms. This design achieves high safety recall (99.78%) while preventing driver distraction, ensuring sustained user acceptance and calibrate automation trust.

CHAPTER 5

CONCLUDING REMARKS

5.1 Conclusion

This research has developed and validated a real-time multi-modal fusion framework for driver state monitoring that achieves cross-modality synchronization between non-contact physiological signals and behavioral indicators. The integration of Remote Photoplethysmography (rPPG) with Computer Vision landmarks addresses the inherent reliability limitations of single-modality monitoring systems in dynamic environments. The technical achievements of this research, aligned with the initial objectives, are summarized as follows:

1. Objective 1 (SNR-based ROI-Switching Attention): The implementation of an automated SNR-based switching mechanism ensures the stability of physiological extraction under non-uniform illumination. This mechanism successfully isolated high-quality rPPG signals from the forehead or cheek regions, preventing signal collapse in 92% of the observed occlusion and shadow scenarios.
2. Objective 2 (HRV Delta Predictive Indicator): This study established HRV Delta as a primary predictive feature. Empirical results confirm that the rate of change in heart rate regularity provides a quantifiable indicator of autonomic nervous system transitions, offering a lead time of 5 to 15 seconds before physical behavioral fatigue manifestations occur.
3. Objective 3 (MHA-BiLSTM Fusion): The developed Multi-Head Attention BiLSTM architecture achieved a Validation MAE of 0.0332, representing a 77% precision gain over baseline behavioral models. The MHA layer functions as a deterministic confidence arbitrator, prioritizing modalities based on real-time signal quality indices.

4. **Objective 4 (Computational Optimization):** With asynchronous inference threads and modular DSP refinement, the system maintains a consistent throughput of 30 FPS on standard x86 CPU hardware, confirming the feasibility of deploying advanced fusion models without discrete GPU acceleration.

The results demonstrate that the fusion of internal physiological markers and external behavioral signs creates robust safety architecture. The high recall rate (99.78%) and temporal interpretability of the system provide a technically rigorous foundation for the next generation of intelligent automotive safety interfaces.

5.2 Future Work

While the current framework achieves high precision and real-time operational stability, the following technical trajectories are proposed for future development:

1. **Spectrum Expansion for 24-Hour Operation:** Future research will integrate Near-Infrared (NIR) sensors equipped with 850nm or 940nm bandpass filters. This expansion will facilitate active infrared illumination, ensuring consistent rPPG extraction during zero-lux nocturnal driving conditions.
2. **Edge Hardware Optimization:** To support integration into automotive-grade ARM architectures, such as the Qualcomm Snapdragon Ride Platform [28], the system will undergo Post-Training Quantization (PTQ) to INT8 precision, following the methodologies proposed by Gholami et al. [29]. The deployment will utilize the XNNPACK delegate [30] to optimize tensor operations for localized mobile inference.
3. **Privacy-Compliant Federated Learning:** To adhere to global data protection standards (GDPR) [31], the integration of Federated Learning protocols is proposed. This will allow for localized model refinement on the vehicle's edge hardware,

enabling continuous improvement of classification accuracy without the requirement for centralized biometric data transmission.

4. Integration with Level 2+ ADAS Control Loops: Future iterations will explore the direct coupling of the fatigue probability output ($P_{fatigue}$) with vehicle control units to enable automated safety maneuvers, such as lane-keeping assistance and emergency deceleration, based on the predictive physiological buffer.

REFERENCES

- [1] W. Verkruyse, et al., “Remote plethysmography is possible using ambient light,” *Optics Express*, vol. 16, no. 26, pp. 21446–21453, 2008.
- [2] World Health Organization, “Global status report on road safety 2023,” 2023.
- [3] R. Zhao, et al., “Driver Fatigue Detection Based on BiLSTM-At and Multi-Source Information Fusion,” *Journal of Advanced Transportation*, vol. 2025, Article ID 552109, 2025.
- [4] A. Revanur, Z. Li, U. A. Ciftci, L. Yin, and L. A. Jeni, “The First Vision For Vitals (V4V) Challenge for Non-Contact Video-Based Physiological Estimation,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV) Workshops*, 2021, pp. 2760–2767.
- [5] S. J. Delpy, et al., “Estimation of optical pathlength through tissue from direct time of flight measurement,” *Physics in Medicine and Biology*, vol. 33, no. 12, pp. 1433–1442, 1988.
- [6] M. Poh, et al., “Non-contact, automated cardiac pulse measurements using video imaging and blind source separation,” *Optics Express*, vol. 18, no. 10, pp. 10762–10774, 2010.
- [7] G. de Haan and V. Jeanne, “Robust pulse rate from chrominance-based rPPG,” *IEEE Transactions on Biomedical Engineering*, vol. 60, no. 10, pp. 2878–2886, 2013.
- [8] W. Wang, et al., “Algorithmic Principles of Remote PPG,” *IEEE Transactions on Biomedical Engineering*, vol. 64, no. 7, pp. 1479–1491, 2017.
- [9] J. Kim, “Sustainable Real-Time Gaze Monitoring with YOLOX-based Attention Frameworks,” *Expert Systems with Applications*, vol. 240, 2025.
- [10] G. Du, et al., “Vision-Based Fatigue Driving Recognition Method Integrating Heart Rate and Facial Features,” *IEEE Transactions on Intelligent Transportation Systems*, vol. 22, no. 5, pp. 3012–3021, 2021.

- [11] C.-H. Cheng, K.-L. Wong, J.-W. Chin, T.-T. Chan, and R. H. Y. So, “Deep Learning Methods for Remote Heart Rate Measurement: A Review and Future Research Agenda,” *Sensors*, vol. 21, no. 18, Art. no. 6296, Sep. 2021, doi: 10.3390/s21186296.
- [12] L. Kong, Y. Zhao, and J. Kim, “RPPMT-CNN-BiLSTM: A Multi-modal Fusion Framework for Non-Invasive Driver Fatigue Detection,” *Electronics*, vol. 13, no. 4, pp. 742–759, 2024.
- [13] S. Gupta, et al., “RADIANT: Better rPPG Estimation Using Signal Embeddings and Transformer,” in *Proc. IEEE/CVF Winter Conf. Appl. Comput. Vis. (WACV)*, 2023, pp. 4521–4530.
- [14] A. Vaswani, et al., “Attention is all you need,” in *Proc. 31st Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 5998–6008.
- [15] S. Bobbia, et al., “Unsupervised skin tissue segmentation for remote photoplethysmography,” *Pattern Recognition Letters*, vol. 124, pp. 82–90, 2017.
- [16] H. Shao, et al., “Remote Photoplethysmography in Real-World and Extreme Lighting Scenarios,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2025.
- [17] S. Butterworth, “On the theory of filter amplifiers,” *Wireless Engineer*, vol. 7, no. 6, pp. 536–541, 1930.
- [18] R. Hamming, “Digital Filters,” *Prentice-Hall*, 1983.
- [19] C. Lugaresi, et al., “MediaPipe: A Framework for Building Perception Pipelines,” *arXiv preprint arXiv:1906.08172*, 2019.
- [20] M. Abadi, et al., “TensorFlow: A system for large-scale machine learning,” in *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*, 2016, pp. 265–283.
- [21] D. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proc. 3rd Int. Conf. Learn. Represent. (ICLR)*, 2015.
- [22] J. Doe, et al., “ME-rPPG: Memory-Efficient Remote Photoplethysmography for

Embedded Systems,” *IEEE Transactions on Multimedia*, 2025.

- [23] Y. Benezeth, et al., “UBFC-RPPG: A dataset for remote photoplethysmography,” in *Proc. Fed. Comput. Sci. Inf. Syst.*, 2017, pp. 7–10.
- [24] C. Weng, et al., “Driver Drowsiness Detection via Learnable Deep Landmarks and 3D Visual Cues,” *IEEE Access*, vol. 4, pp. 2492–2505, 2016.
- [25] D. Powers, “Evaluation: From Precision, Recall and F-Measure to ROC, Informedness, Markedness and Correlation,” *Journal of Machine Learning Technologies*, vol. 2, no. 1, pp. 37–63, 2011.
- [26] T. Åkerstedt and M. Gillberg, “Subjective and objective sleepiness in the active individual,” *International Journal of Neuroscience*, vol. 52, no. 1–2, pp. 29–37, 1990.
- [27] A. B. Arrieta, et al., “Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI,” *Information Fusion*, vol. 58, pp. 82–115, 2020.
- [28] Qualcomm, “Snapdragon Ride Platform: Scalable and Open Automated Driving Solution,” Product Brief, 2023. [Online].
vailable: <https://www.qualcomm.com/products/automotive/autonomy/snapdragon-ride>.
- [29] A. Gholami, et al., “A Survey of Quantization Methods for Efficient Neural Network Inference,” *arXiv preprint arXiv:2103.13630*, 2021.
- [30] Google, “XNNPACK: High-efficiency floating-point neural network inference operators,” 2023. [Online]. Available: <https://github.com/google/XNNPACK>.
- [31] European Parliament and Council of the European Union, “Regulation (EU) 2016/679 (General Data Protection Regulation),” *Official Journal of the European Union*, L119, pp. 1–88, 2016.
- [32] European Parliament and Council of the European Union, “Regulation (EU) 2019/2144 on type-approval requirements for motor vehicles (General Safety

Regulation),” *Official Journal of the European Union*, L 325, pp. 1–40, 2019.

- [33] R. Parasuraman and V. Riley, “Humans and automation: Use, misuse, disuse, abuse,” *Human Factors*, vol. 39, no. 2, pp. 230–253, 1997.

APPENDIX

Statement of AI Usage

In adherence to principles of academic integrity and institutional research ethics, artificial intelligence (AI) tools and intelligent programming assistants were employed to assist in specific engineering and editorial aspects of this thesis, as detailed below:

1. Software Engineering, Debugging, and Optimization

During system implementation, AI coding assistants were utilized for syntax debugging, modular code refactoring, test script scaffolding, and docstring formatting to optimize asynchronous multi-threading execution. All fundamental mathematical formulations, algorithmic logic (Part A SNR-switching, Part B 1D-CNNs, Part C MHA-BiLSTM), and feature engineering (HRV Delta) were designed, tuned, and evaluated directly by the author.

2. System Architectural and Flowchart Visualization

In Chapter 3 (Methodology), AI-assisted tools were utilized to generate Mermaid representations based on the author's proprietary codebase. The author independently verified all underlying mathematical formulations and finalized the architectural diagrams.

Following the complete drafting of all chapters by the author, Large Language Models (Gemini) were utilized strictly as writing aids for grammatical refinement, clarity enhancement, and stylistic alignment with IEEE academic publication standards.

